

SEC204

Computer architectures and low level
programming
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Email: v.kelefouras@plymouth.ac.uk

Website: <https://www.plymouth.ac.uk/staff/vasilios-kelefouras>

Outline

2

- Different ways of writing assembly code
- Using intrinsic functions in C/C++
- Writing C/C++ programs using Intel SSE intrinsics
- Writing C/C++ programs using Intel AVX intrinsics

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Different ways of writing assembly

3

1. **Writing an entire function in assembly**
 2. **Using inline assembly in C/C++**
 3. **Using intrinsic functions in C/C++**
 - highly recommended - much easier and safer
 - All the compilers support intrinsic functions
 - An intrinsic function is equivalent to an assembly instruction
 - **Mixes the good things of C++** (development time, portability, maintainability etc) **with the good things of assembly** (execution time)
- C and C++ are the most easily combined languages with assembly code

Different ways of writing assembly

Using intrinsic functions in C/C++

4

□ Main advantages

- ▣ Classes, if conditions, loops and functions are very easy to implement
- ▣ Portability to almost all x86 architectures
- ▣ Compatibility with different compilers

□ Main disadvantages

- ▣ Not all assembly instructions have intrinsic function equivalents
- ▣ Unskilled use of intrinsic functions can make the code less efficient than simple C++ code

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Using intrinsic functions in C/C++

5

- **For the rest of this lecture, you will be learning how to use intrinsic functions in C/C++**

Assignment Project Exam Help

- Normally, “90% of a program’s execution time is spent in executing 10% of the code” - <https://powcoder.com>
 - ▣ What programmers normally do to improve performance is to analyze the code and find the computationally intensive functions
 - Then optimize those instead of the whole program
 - This saves time and money
 - ▣ **Rewriting loop kernels in C++ using SIMD intrinsics is an excellent choice**
 - Compilers vectorize the code (not always) but manually using SIMD intrinsics can really boost performance

Single Instruction Multiple Data (SIMD) – Vectorization

6

Assignment Project Exam Help

```
for (i = 0; i < MAT_SIZE; i++)  
    c[i] = a[i] + b[i];
```

<https://powcoder.com>

Add WhatsApp powcoder



Vectorization on Arm Cortex series

NEON technology

7

- Arm Neon technology is an advanced SIMD architecture extension for the Arm Cortex-A series and Cortex-R52 processors
 - ▣ 128-bit wide
 - ▣ They are widely used in embedded systems
- Neon instructions allow up to:
 - ▣ 16x8-bit, 8x16-bit, 4x32-bit, 2x64-bit integer operations
 - ▣ 8x16-bit, 4x32-bit, 2x64-bit floating-point operations

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Vectorization on Intel Processors

- **Intel MMX technology** (old – limited usage nowadays)
 - 8 mmx registers of 64 bit
 - extension of the floating point registers
 - can be handled as 8 8-bit, 4 16-bit, 2 32-bit and 1 64-bit, operations
 - An entire L1 cache line is loaded to the RF in 1-3 cycles
- **Intel SSE technology**
 - 8/16 xmm registers of 128 bit (32-bit architectures support 8 registers only)
 - Can be handled from 16 8-bit to 1 128-bit operations
 - An entire L1 cache line is loaded to the RF in 1-3 cycles
- **Intel AVX technology**
 - 8/16 ymm registers of 256 bit (32-bit architectures support 8 registers only)
 - Can be handled from 32 8-bit to 1 256-bit operations
- **Intel AVX-512 technology**
 - 32 ZMM 512-bit registers

Vectorization on Intel Processors (2)

9


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

SSE and AVX-128 types

AVX-256 types



Vectorization on Intel Processors (3)

- The developer can use either SSE or AVX or both
 - ▣ AVX instructions improve throughput
 - ▣ SSE instructions are preferred for less data parallel algorithms
- Vector instructions work only for data that they are written in consecutive main memory addresses
- Aligned load/store instructions are faster than the no aligned ones.
- memory and arithmetical instructions are executed in parallel

- **All the Intel intrinsics can be found here :**

<https://software.intel.com/sites/landingpage/IntrinsicsGuide/#>

Basic SSE Instructions (1)

11

- `__m128 _mm_load_ps (float * p)` – Loads four SP FP values. The address must be 16-byte-aligned
- `__m128 _mm_loadu_ps (float * p)` – Loads four SP FP values. The address need not be 16-byte-aligned

Assignment Project Exam Help

L1

A[0]	A[1]	A[2]	A[3]
A[4]	A[5]	A[6]	A[7]
....			

<https://powcoder.com>
Aligned load

L1

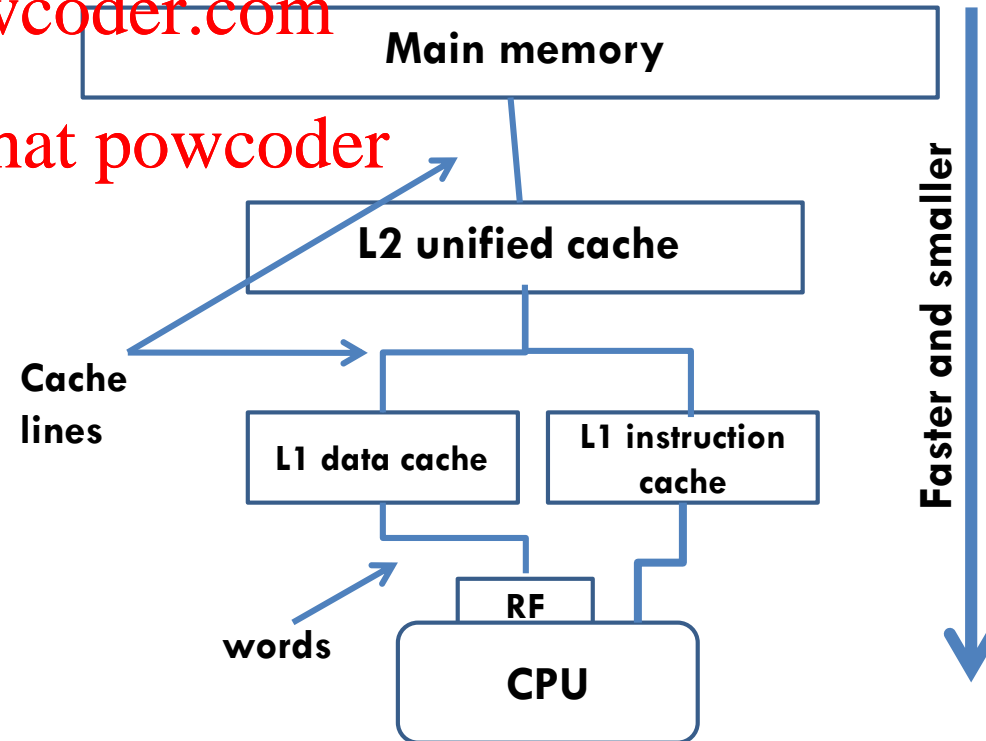
A[0]	A[1]	A[2]	A[3]
A[4]	A[5]	A[6]	A[7]
....			

Misaligned load

L1

A[1]	A[2]	A[3]	A[4]
A[5]	A[6]	A[7]	A[8]
....			

Misaligned load



Basic SSE Instructions (2)

12

- `__m128 _mm_load_ps(float * p)` – Loads four SP FP values. The address must be 16-byte-aligned
- `__m128 _mm_loadu_ps(float * p)` – Loads four SP FP values. The address need not be 16-byte-aligned

Assignment Project Exam Help

`float A[N] __attribute__((aligned(16)));`

<https://powcoder.com>
Aligned load

Add WeChat powcoder

Main Memory



Modulo (Address ,16)=0

L1

A[0]	A[1]	A[2]	A[3]
A[4]	A[5]	A[6]	A[7]
....			

L1

A[0]	A[1]	A[2]	A[3]
A[4]	A[5]	A[6]	A[7]
....			

Misaligned load

L1

A[0]	A[1]	A[2]	A[3]
A[4]	A[5]	A[6]	A[7]
....			


Misaligned load

Basic SSE Instructions (3)


13

- `__m128 _mm_store_ps(float * p)` – Stores four SP FP values. The address must be 16-byte-aligned
- `__m128 _mm_storeu_ps(float * p)` – Stores four SP FP values. The address need not be 16-byte-aligned
- `__m128 _mm_mul_ps(__m128 a, __m128 b)` – Multiplies the four SP FP values of a and b
- `__m128 _mm_mul_ss(__m128 a, __m128 b)` – Multiplies the lower SP FP values of a and b; the upper 3 SP FP values are passed through from a.

`XMM1 = _mm_mul_ss(XMM1, XMM0)`



`XMM1 = _mm_mul_ps(XMM1, XMM0)`



Basic SSE Instructions (4)

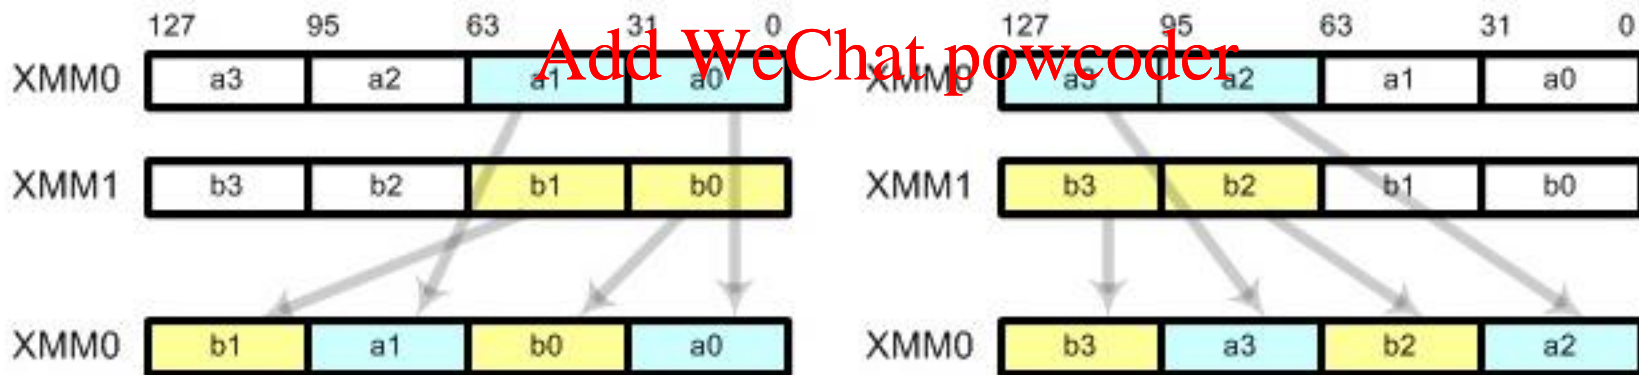
14

- `__m128 _mm_unpackhi_ps (__m128 a, __m128 b)` - Selects and interleaves the upper two SP FP values from a and b.
- `__m128 _mm_unpacklo_ps (__m128 a, __m128 b)` - Selects and interleaves the lower two SP FP values from a and b.

Assignment Project Exam Help

<https://powcoder.com>


`XMM0=_mm_unpacklo_ps (XMM0, XMM1)` `XMM0=_mm_unpackhi_ps (XMM0, XMM1)`



Basic SSE Instructions (5)

15

- `__m128 _mm_hadd_ps (__m128 a, __m128 b)` - Adds adjacent vector elements



- `void _mm_store_ss (float * p, __m128 a)` - Stores the lower SP FP value

Case Study

MVM using SSE technology

16

```
float A[N][N];  
float X[N], Y[N];  
int i,j;
```

Assignment Project Exam Help

```
for (i=0; i<N; i++){
```

```
float A[N][N];  
float X[N], Y[N];  
int i,j;
```

```
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    Y[i] += A[i][j] * X[j];
```

<https://powcoder.com>

```
num3=_mm_setzero_ps();
```

Add WeChat powcoder

```
for (j=0; j<N; j+=4){
```

```
  num0=_mm_load_ps( &A[i][j] );
```

```
  num1=_mm_load_ps(X + j );
```

```
  num3=_mm_fmadd_ps(num0,num1,num3);
```

```
}
```

```
num4=_mm_hadd_ps(num3, num3);
```

```
num4=_mm_hadd_ps(num4, num4);
```

```
_mm_store_ss((float *)Y+i, num4);
```

```
}
```


MVM with SSE (2)

```
for (i=0; i!=N; i++){
```

```
num3=_mm_setzero_ps();
```



```
for (j=0; j!=N; j+=4){
```

```
num0=_mm_load_ps( &A[i][j] );
```

```
num1=_mm_load_ps(X + j );
```


```
num3=_mm_fmadd_ps(num0,num1,num3);
```



```
num3=_mm_hadd_ps(num3, num3);
```

```
num3=_mm_hadd_ps(num3, num3);
```

```
_mm_store_ss((float *)Y+i, num3);
```

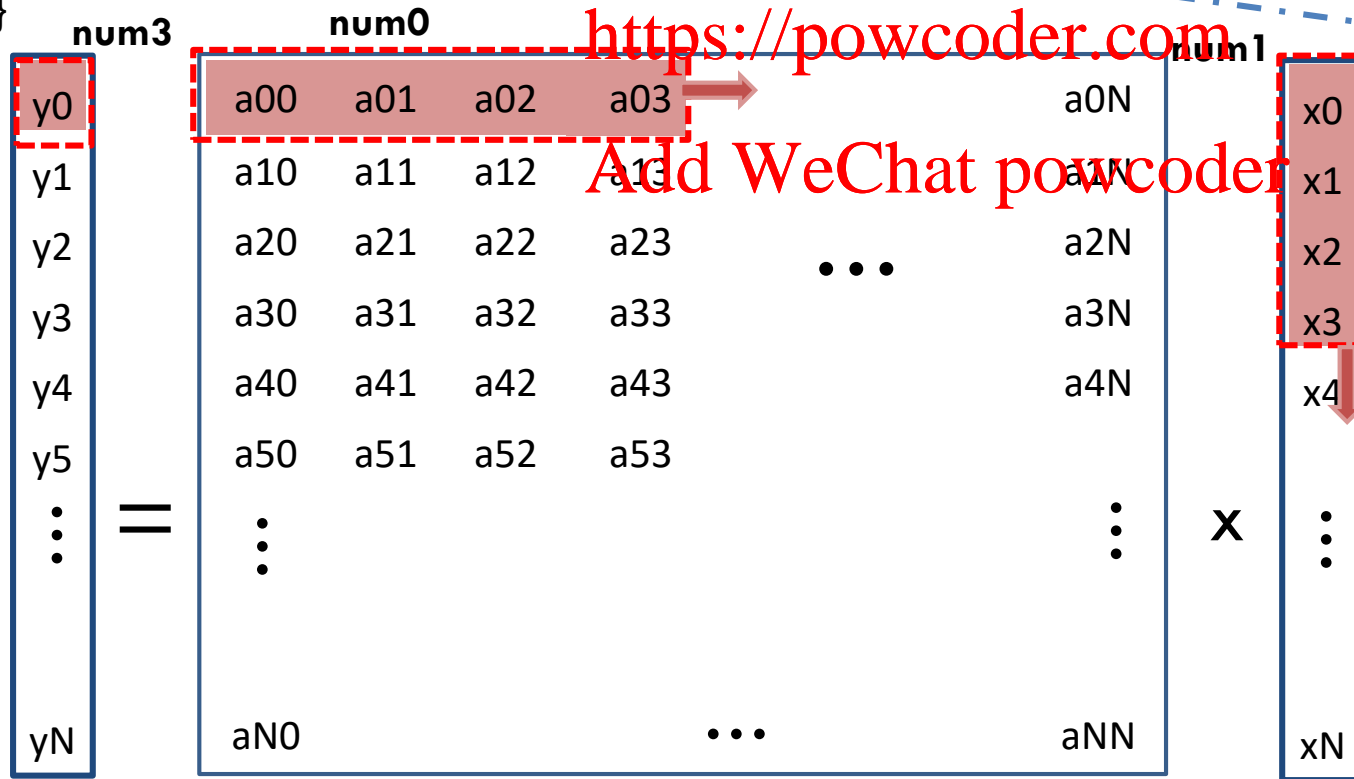


Assignment Project Exam Help

<https://powcoder.com>

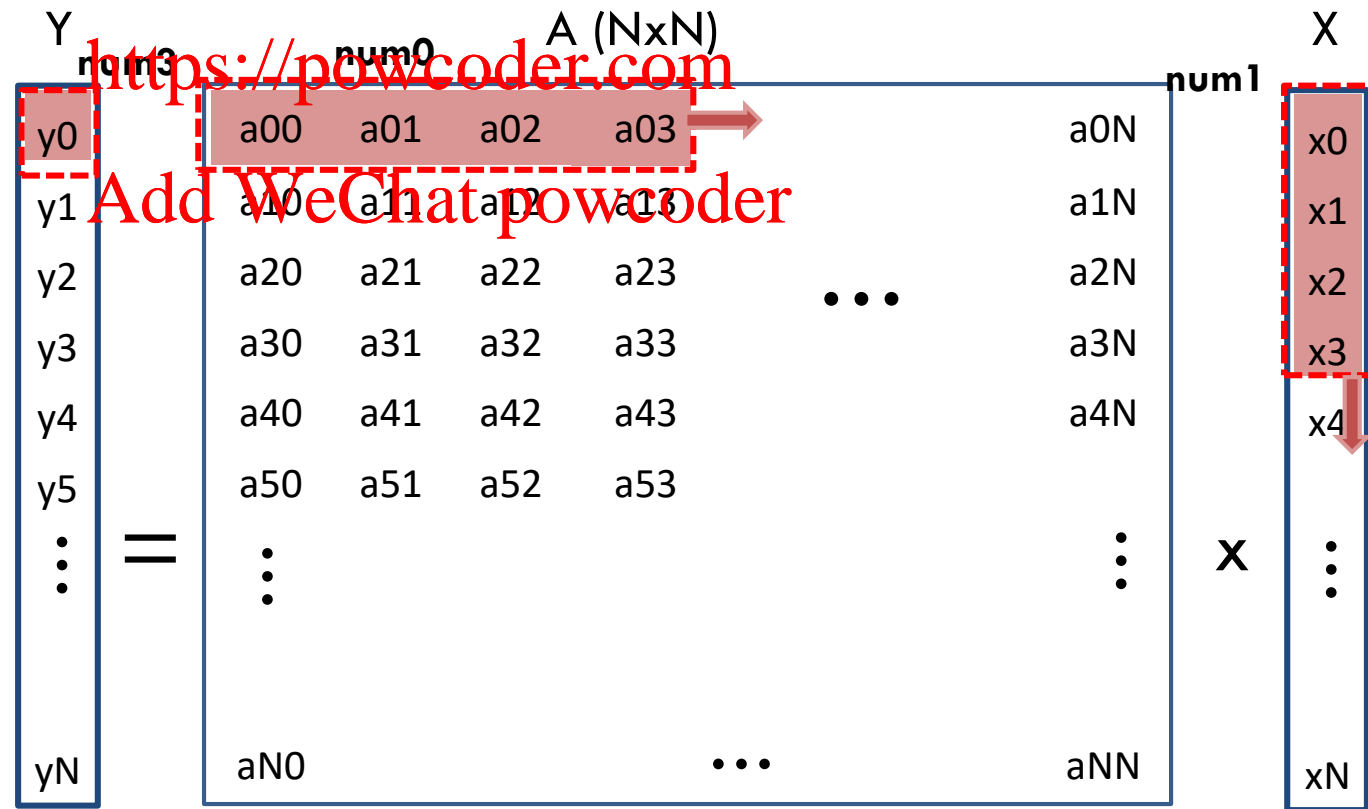
Add WeChat powcoder

This part of code adds the four values of num3 and stores the result into Y[i]



```
....
}
num3=_mm_hadd_ps(num3, num3);
num3=_mm_hadd_ps(num3, num3);
_mm_store_ss((float *)Y+i, num3);
}
```

- After j loop finishes its execution, num3 contains the output data of Y[i]
- num3=[ya, yb, yc, yd] where $Y[i]=ya+yb+yc+yd$**
- after the 1st hadd \rightarrow num3=[ya+yb, yc+yd, ya+yb, yc+yd]
- after the 2nd hadd \rightarrow num3=[ya+yb+yc+yd, ya+yb+yc+yd, ya+yb+yc+yd, ya+yb+yc+yd]



Case Study

MVM using AVX technology

19

```
float A[N][N];  
float X[N], Y[N];  
int i,j;  
  
for (i=0; i<N; i++)  
    for (j=0; j<N; j++)  
        Y[i] += A[i][j] * X[j];
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

```
for (i=0; i!=N; i++)  
    for (j=0; j!=N; j++) {  
        ymm0 = _mm256_setzero_ps();  
        for (k=0; k!=N; k+=8) {  
            ymm1 = _mm256_load_ps( A + N*i + k);  
            ymm2 = _mm256_load_ps( Btrans + N*j + k);  
            ymm0 = _mm256_fmadd_ps(ymm1, ymm2, ymm0);  
        }  
  
        ymm2 = _mm256_permute2f128_ps(ymm0, ymm0, 1);  
        ymm0 = _mm256_add_ps(ymm0, ymm2);  
        ymm0 = _mm256_hadd_ps(ymm0, ymm0);  
        ymm0 = _mm256_hadd_ps(ymm0, ymm0);  
        _mm_store_ss((float *) C + N*i + j,  
                    _mm256_extractf128_ps(ymm0, 0));  
    }  
}
```

```
for (i=0; i < n; i++) {
    if ( x[i] > 2 || x[i] < -2 )
        a[i]+=x[i]; }

```

What about if-conditions on SSE ?



```
const __m128 P2f = _mm_set1_ps(2.0f);
const __m128 M2f = _mm_set1_ps(-2.0f);
for (int i = 0; i < n; i += 4)
{
    __m128 xv = _mm_load_ps(x + i);
    __m128 av = _mm_load_ps(a + i);

    __m128 c1v = _mm_cmpgt_ps(xv, P2f);
    __m128 c2v = _mm_cmplt_ps(xv, M2f);

    __m128 cv = _mm_or_ps(c1v, c2v);

    xv = _mm_and_ps(xv, cv);
    av = _mm_add_ps(av, xv);

    _mm_store_ps(a + i, av);
}

```

2	2	2	2
---	---	---	---

-2	-2	-2	-2
----	----	----	----

5	-3	0	1
---	----	---	---

a[i]	a[i+1]	a[i+2]	a[i+3]
------	--------	--------	--------

1	0	0	0
---	---	---	---

0	1	0	0
---	---	---	---

1	1	0	0
---	---	---	---

x[i]	x[i+1]	0	0
------	--------	---	---

a[i] +	a[i+1] +	a[i+2] +	a[i+3] +
x[i]	x[i+1]	0	0

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder